

Final Project Report

DevOps Pipeline for a Multi-Service Application

Giorgi Gagnidze

June 2025

Table of Contents

1. Introduction.....	2
2. General Project Structure.....	3
3. Application Overview	4
4. Containerize Services.....	6
5. Compose the Stack.....	9
6. Monitoring & Visualization.....	10
7. Security	18
8. Incident Simulation + Post-Mortem.....	22
9. Automation (Ansible, Terraform, GitHub Actions)	22

1. Introduction

This project demonstrates the design and implementation of a modern DevOps pipeline for a multi-service application. The main objective is to containerize and orchestrate multiple services using Docker and Docker Compose, while integrating monitoring, security, and automation practices throughout the application lifecycle.

Key components of the project include:

- **Containerization:** Both frontend and backend services are packaged into independent containers with custom Dockerfiles and served by NGINX.
- **Orchestration:** Docker Compose is used to define and run the entire stack, ensuring inter-service communication and simplified container management.
- **Monitoring & Visualization:** Prometheus collects real-time metrics from the backend service, and Grafana provides interactive dashboards to visualize application health and performance, alert manager is responsible for sending alert emails to maintainers in case of a failure.
- **Stress Testing:** Stress testing with Gatling, more than 1000 concurrent requests for CRUD operations, and on-purpose, increased CPU & memory usage within each request to better visualize utilization metrics data in Grafana.
- **Security:** The security of all Docker images is assessed using Trivy. Sensitive data is managed securely using .env files and GitHub Actions secrets, as well as Ansible is responsible for securely uploading .env files on the VM.
- **Automation:** Ansible is used to automate deployment and provisioning tasks (updating Linux repositories, creating folders, etc...) reducing manual intervention and ensuring smooth deployment process to have the infrastructure ready to go.

- **Infrastructure as Code (IaC):** Terraform provisions the cloud infrastructure, in this case GCP virtual machine and manages it automatically, ensuring reproducibility and ease of scaling.
- **Version Control:** The entire project is tracked with Git and hosted on GitHub, which also provides GitHub actions, enabling effective collaboration, transparency and automation. Deploy is done on a Terraform provided machine

2. General Project Structure

devops-final-project/

```

├── .github/           # GitHub actions
├── backend/           # Spring Boot backend service
│   ├── build/
│   ├── checkstyle.xml
│   ├── Dockerfile
├── frontend/          # React frontend service
│   ├── build/
│   ├── public/
│   ├── src/
│   ├── .env.development
│   ├── Dockerfile
│   ├── package.json
│   ├── package-lock.json
│   └── tsconfig.json
├── infra/             # Infrastructure and deployment
│   ├── ansible/
│   ├── backend/      # ENV
│   ├── frontend/     # ENV
│   ├── gatling/       # Gatling results and simulation
│   └── monitoring/    # Grafana, Prometheus and Alert manager

```

- | |— nginx/ # Folder containing nginx.conf
- | |— terraform/ # terraform files provisioning GCP Instance
- | |— compose.yml # Docker compose file
- | |— scan-all.sh # Trivy script that scans all Docker images

3. Application Overview

The application is a full-stack, cloud-deployed system, consisting of a Spring Boot backend REST API (with Swagger) and a ReactJS frontend served by NGINX. The backend exposes RESTful APIs for creating TODO-tasks which are updateable, deletable and savable.

The Backend comes with H2 embedded database and exposes health check endpoints for Prometheus and Grafana. It also comes with checkstyle and JUnit integration tests.

```
spring.datasource.url=${SPRING_DATASOURCE_URL:jdbc:h2:mem:testdb}
spring.datasource.username=${SPRING_DATASOURCE_USERNAME:sa}
spring.datasource.password=${SPRING_DATASOURCE_PASSWORD:}
spring.datasource.driverClassName=org.h2.Driver

management.endpoints.web.exposure.include=*
management.endpoint.prometheus.access=unrestricted
management.prometheus.metrics.export.enabled=true
```

Architecture Overview

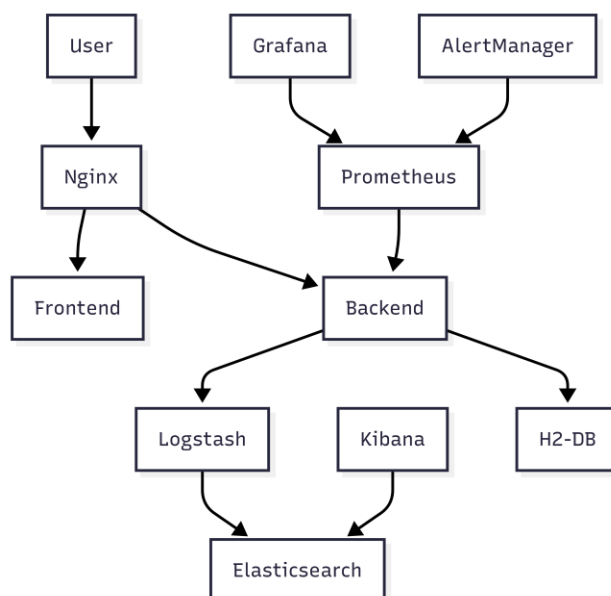
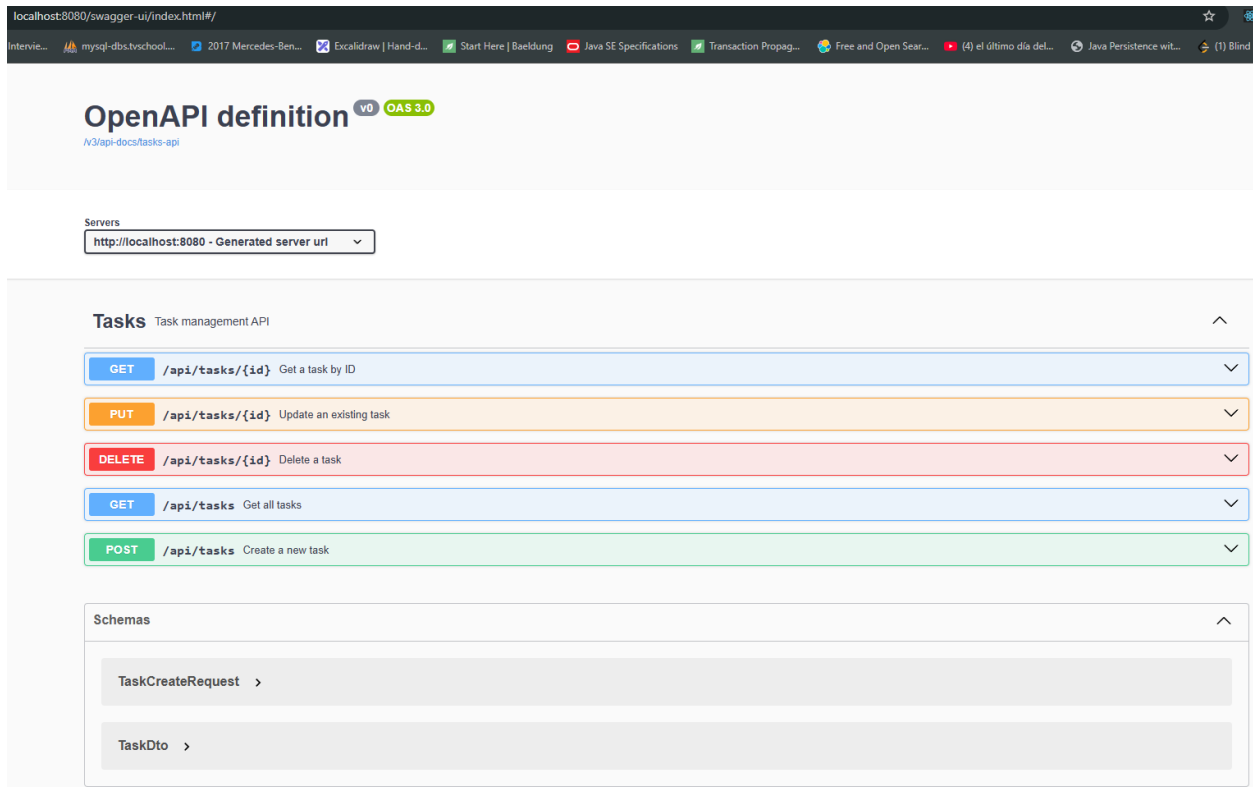


Image above displays the general flow, how the user is served by NGINX and the backend communication with Alerting and Monitoring services.

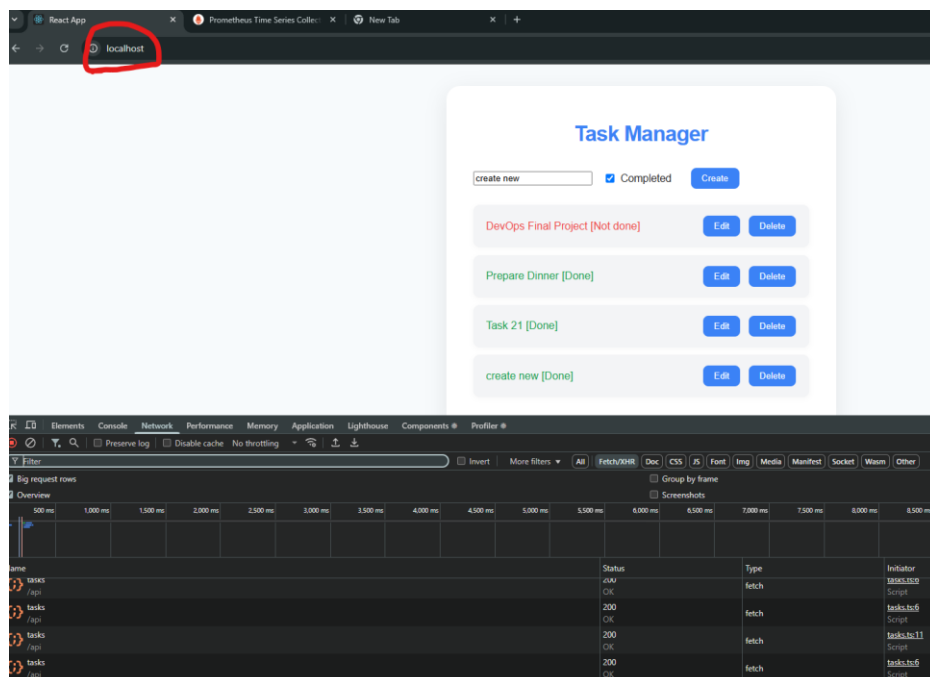
Demo

1) API:



2) UI served by NGINX:

The user can create, update, or delete tasks from the dashboard



4. Containerize Services

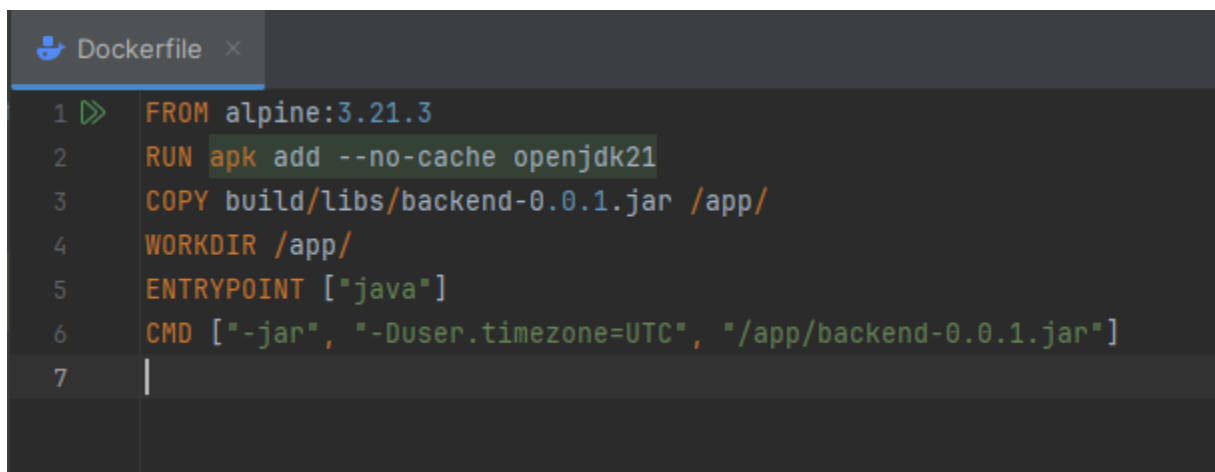
The task was to create at least two services (backend, frontend) and ensure services ran independently in the containers.

Under my root directory I've created 3 folders

- Backend – containing backend service with Dockerfile
- Frontend – containing front end service with its own Dockerfile as well
- Infra – containing all of the infrastructure related stuff, as well as Compose.yml file

Backend

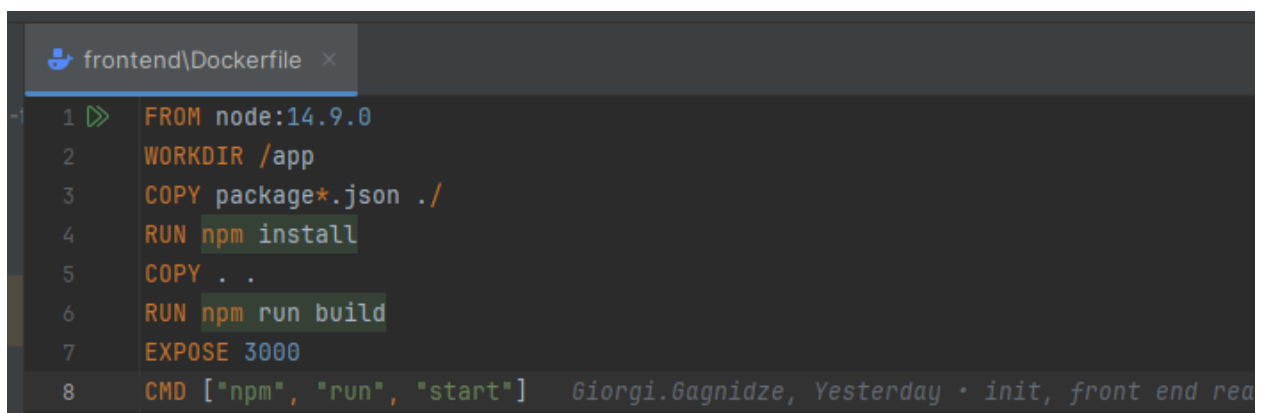
This is the Dockerfile for my backend service, copying generated jar from the build folder to the app directory and providing a command for running the app

A screenshot of a code editor showing a Dockerfile. The file is named 'Dockerfile' and is located in a folder. The code consists of seven lines: 1. 'FROM alpine:3.21.3', 2. 'RUN apk add --no-cache openjdk21', 3. 'COPY build/libs/backend-0.0.1.jar /app/', 4. 'WORKDIR /app/', 5. 'ENTRYPOINT ["java"]', 6. 'CMD ["-jar", "-Duser.timezone=UTC", "/app/backend-0.0.1.jar"]', and 7. an empty line. The code is syntax-highlighted with various colors.

```
1 FROM alpine:3.21.3
2 RUN apk add --no-cache openjdk21
3 COPY build/libs/backend-0.0.1.jar /app/
4 WORKDIR /app/
5 ENTRYPOINT ["java"]
6 CMD ["-jar", "-Duser.timezone=UTC", "/app/backend-0.0.1.jar"]
7
```

Frontend

This is the Dockerfile for my frontend service

A screenshot of a code editor showing a Dockerfile. The file is named 'frontend\Dockerfile'. The code consists of eight lines: 1. 'FROM node:14.9.0', 2. 'WORKDIR /app', 3. 'COPY package*.json ./', 4. 'RUN npm install', 5. 'COPY . .', 6. 'RUN npm run build', 7. 'EXPOSE 3000', and 8. 'CMD ["npm", "run", "start"]'. The code is syntax-highlighted with various colors.

```
1 FROM node:14.9.0
2 WORKDIR /app
3 COPY package*.json ./
4 RUN npm install
5 COPY . .
6 RUN npm run build
7 EXPOSE 3000
8 CMD ["npm", "run", "start"]
```

Building Containers

Here's my compose.yml file snippet containing backend and front end images

```
compose.yml x
1  name: devops-final-prj
2
3  services:
4    backend:
5      container_name: backend
6      restart: always
7      build:
8        context: ../backend
9        dockerfile: Dockerfile
10     ports:
11       - "8080:8080"
12     depends_on:
13       - grafana
14       - prometheus
15       - alertmanager
16       - gatling
17     env_file:
18       - backend/.env.development
19
20   frontend:
21     container_name: frontend
22     build:
23       context: ../frontend
24       dockerfile: Dockerfile
25     env_file:
26       - frontend/.env.development
27     expose:
28       - "5000"
29     depends_on:
30       - nginx
```

Building backend with docker compose build --no-cache backend

```
PS C:\Users\giorg\Desktop\devops-final-project\infra> docker compose build --no-cache backend
[+] Building 37.2s (10/10) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 32B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/alpine:3.21.3
=> [auth] library/alpine:pull token for registry-1.docker.io
=> CACHED [1/4] FROM docker.io/library/alpine:3.21.3@sha256:a8560b36e8b8210634f77d9f7f9efd7ffa463e380b75e2e74aff4511df3ef88c
=> [internal] load build context
```

Building frontend with docker compose build --no-cache frontend

```
Terminal  Local  ×  Ubuntu-20.04  ×  Ubuntu-20.04 (2)  ×  +  ∨
PS C:\Users\giorg\Desktop\devops-final-project\infra> docker compose build --no-cache frontend
[+] Building 92.2s (12/12) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 32B
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [internal] load metadata for docker.io/library/node:14.9.0
=> [auth] library/node:pull token for registry-1.docker.io
=> [1/6] FROM docker.io/library/node:14.9.0@sha256:ce506ed8986a0c8a364757771679706ebd129fa466165fcc6e2c7dc449a0baac
=> [internal] load build context
```

<input type="checkbox"/>	Name	Image	Status	Port(s)	Started	Actions
<input type="checkbox"/>	devops-final-prj	-	Running (10/1)			
<input type="checkbox"/>	alertmanager e9f7b30bc1c	prom/alertmanager:v0.26.0	Running	9093:9093	17 seconds ago	
<input type="checkbox"/>	elasticsearch 529250e19ae	docker.elastic.co/elasticsearch/elasticsearch:8.13.4	Running	9200:9200	17 seconds ago	
<input type="checkbox"/>	nginx e506b900e698	nginx:1.25-alpine	Running	80:80	17 seconds ago	
<input type="checkbox"/>	grafana 65cc67bbe6d	grafana/grafana:9.5.2	Running	3001:3000	18 seconds ago	
<input type="checkbox"/>	gatling d8b6ac37cd3df	demich/gatling	Running		18 seconds ago	
<input type="checkbox"/>	prometheus ce3ba9315431	prom/prometheus:v2.44.0	Running	9090:9090	17 seconds ago	
<input type="checkbox"/>	kibana 468d61da886b	docker.elastic.co/kibana/kibana:8.13.4	Running	5601:5601	16 seconds ago	
<input type="checkbox"/>	frontend 3848161a1523	devops-final-prj-frontend	Running		16 seconds ago	
<input type="checkbox"/>	backend 43385f479533	devops-final-prj-backend	Running	8080:8080	16 seconds ago	
<input type="checkbox"/>	logstash 542848842ab4	docker.elastic.co/logstash/logstash:8.13.4	Running	5555:5555 9600:9600	16 seconds ago	

NGINX Configuration

```
nginx.conf
1 worker_processes 1;
2 events { worker_connections 1024; }
3
4 http {
5     include mime.types;
6     default_type application/octet-stream;
7
8     server {
9         listen 80;
10
11         location / {
12             proxy_pass http://frontend:3000/;
13             proxy_set_header Host $host;
14             proxy_set_header X-Real-IP $remote_addr;
15         }
16
17         location /api/ {
18             proxy_pass http://backend:8080/api/;
19             proxy_set_header Host $host;
20             proxy_set_header X-Real-IP $remote_addr;
21         }
22     }
23 }
24
```

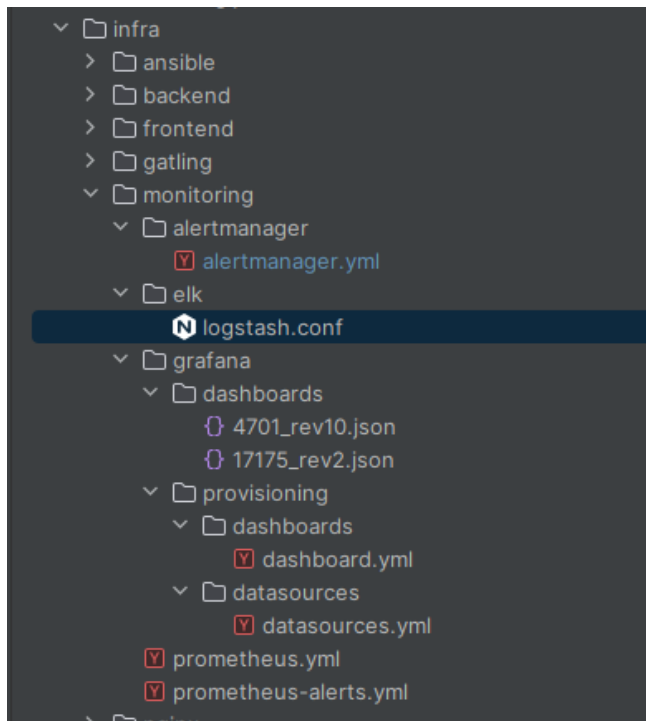
```
nginx:
  image: nginx:1.25-alpine
  container_name: nginx
  ports:
    - "80:80"
  volumes:
    - ./nginx/nginx.conf:/etc/nginx/nginx.conf:ro
```

5. Compose the Stack

```
Terminal Local x Ubuntu-20.04 x Ubuntu-20.04 (2) x + v
PS C:\Users\giorg\Desktop\devops-final-project\infra> docker compose up -d
[*] Running 11/11
- Network devops-final-prj_default Created 0.0s
- Container prometheus Started 2.2s
- Container gatling Started 1.1s
- Container grafana Started 1.4s
- Container elasticsearch Started 2.0s
- Container nginx Started 2.3s
- Container alertmanager Started 2.1s
- Container logstash Started 3.2s
- Container kibana Started 3.0s
- Container frontend Started 3.5s
- Container backend Started 3.5s
PS C:\Users\giorg\Desktop\devops-final-project\infra>
```

By default, all services defined in Compose.yml file are automatically connected to default network, if not defined custom one.

To see that the communication works between services, we can check logs, for example logstash, we see application logs coming from the backend service and connection to the elastic:



The volumes are mounted to the corresponding directories in the project.

As seen here, I have two dashboards for Grafana, one visualizing Spring Boot metrics, the other visualizes JVM Specific metrics aka Micrometer.

Dashboards file tells Grafana which dashboards to use that are provided, datasources define and map prometheus variable as a source of data for the given dashboard templates.

Alertmanager yaml file contains smtp secrets needed for sending emails to the users.

```

1  global:
2    scrape_interval: 15s
3
4  alerting:
5    alertmanagers:
6      - static_configs:
7        - targets:
8          - 'alertmanager:9093'
9
10   rule_files:
11     - 'prometheus-alerts.yml'
12
13   scrape_configs:
14     - job_name: 'backend-metrics'
15       metrics_path: '/actuator/prometheus'
16       scrape_interval: 5s
17       static_configs:
18         - targets: [ 'host.docker.internal:8080' ]
19         labels:
20           application: 'backend-metrics'

```

```

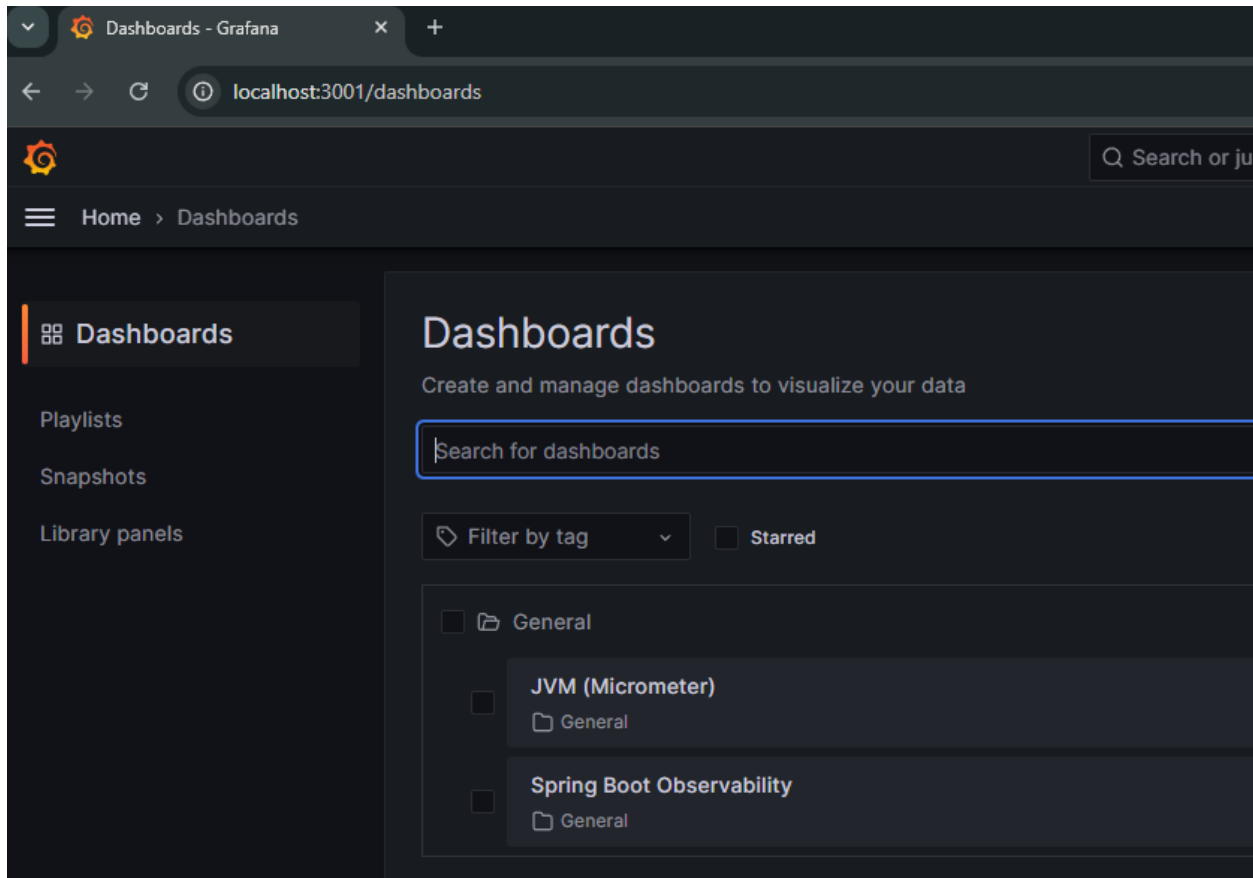
1  groups:
2    - name: service-alerts
3      rules:
4        - alert: BackendDown
5          expr: up{job="backend-metrics"} == 0
6          for: 30s
7          labels:
8            severity: critical
9          annotations:
10            summary: "Backend API is down"
11            description: "The Spring Boot backend (job=
12
13        - alert: HighBackendCPU
14          expr: process_cpu_seconds_total{job="backend-
15          for: 1m
16          labels:
17            severity: warning
18          annotations:
19            summary: "Backend CPU usage is high"
20            description: "Backend process CPU usage has

```

These two files configure Prometheus to scrape the backend health from the actuator endpoint as well as configure communication between Prometheus and alert manager. The file on the right defines two tasks or the rulesets according to which alerts are sent in case of an expression logic is satisfied, for example when the backend service is down for more than 30 seconds and etc.

Visualization

We see the dashboards we defined earlier: Micrometer and Spring Boot related

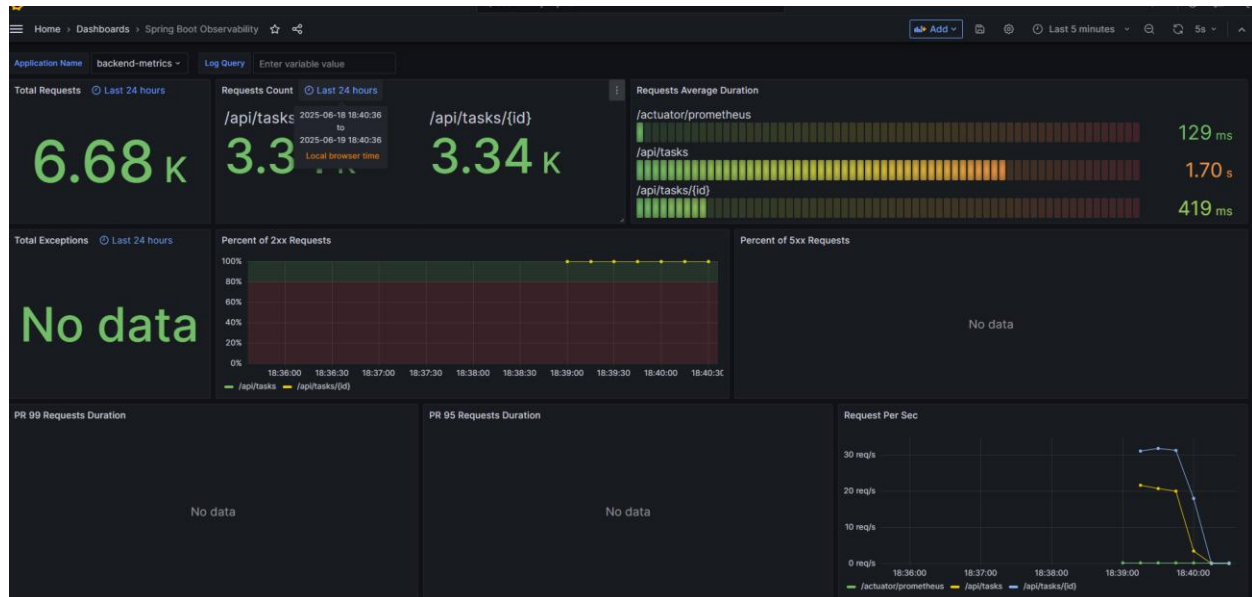


Now, let's observe how the metrics are going, considering Gatling is running after starting the Docker compose entirely.

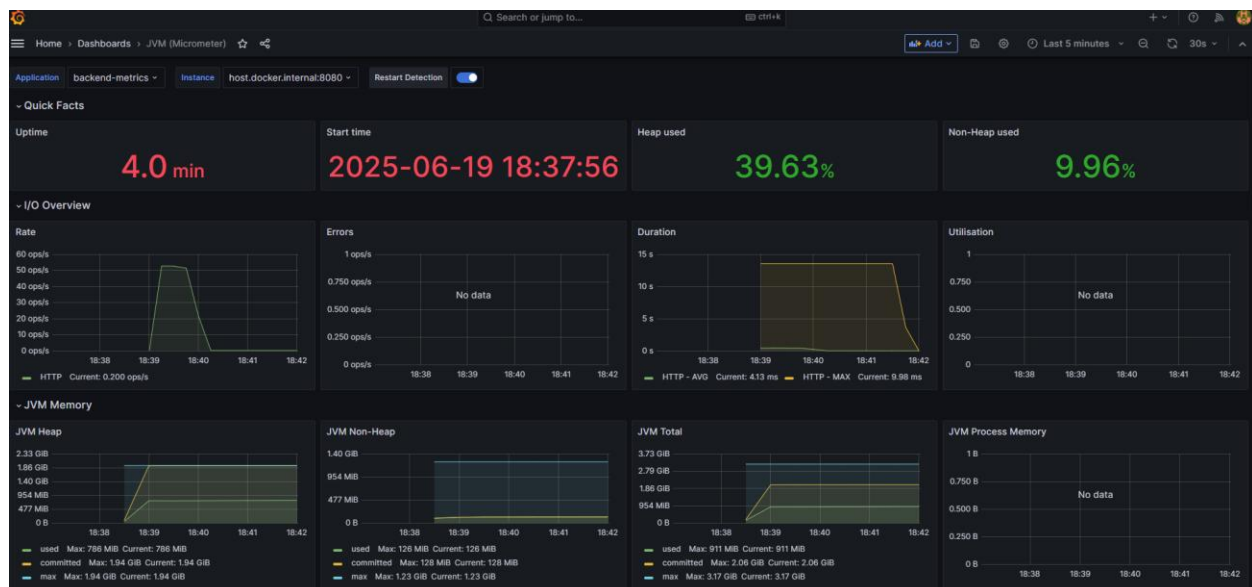
For the first test, I set 2500 concurrent requests for CRUD on the controller, seems its not able to handle it all together, considering there's timeout exceptions in Gatling logs

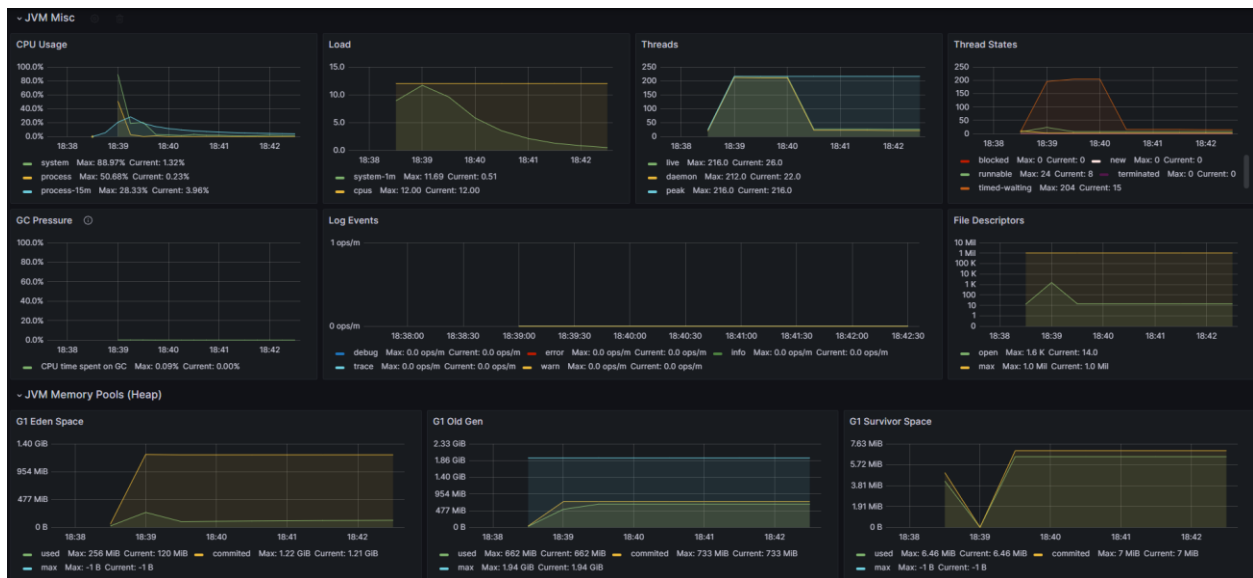
```
at java.lang.Thread.run(Thread.java:748)
14:38:44.396 [WARN ] i.g.h.e.r.DefaultStateProcessor - Request 'Create Task' failed for user 2425: i.n.c.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
14:38:44.397 [WARN ] i.g.h.e.GatlingHttpListener - Request 'Create Task' failed for user 2353
io.netty.channel.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
at io.netty.channel.nio.AbstractNioChannel$AbstractNioUnsafe$1.run(AbstractNioChannel.java:263)
at io.netty.util.concurrent.PromiseTask$RunnableAdapter.call(PromiseTask.java:38)
at io.netty.util.concurrent.ScheduledFutureTask.run(ScheduledFutureTask.java:127)
at io.netty.util.concurrent.AbstractEventExecutor.safeExecute(AbstractEventExecutor.java:163)
at io.netty.util.concurrent.SingleThreadEventExecutor.runAllTasks(SingleThreadEventExecutor.java:416)
at io.netty.channel.nio.NioEventLoop.run(NioEventLoop.java:515)
at io.netty.util.concurrent.SingleThreadEventExecutor$5.run(SingleThreadEventExecutor.java:918)
at io.netty.util.internal.ThreadExecutorMap$2.run(ThreadExecutorMap.java:74)
at io.netty.util.concurrent.FastThreadLocalRunnable.run(FastThreadLocalRunnable.java:30)
at java.lang.Thread.run(Thread.java:748)
14:38:44.397 [WARN ] i.g.h.e.r.DefaultStateProcessor - Request 'Create Task' failed for user 2353: i.n.c.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
14:38:44.397 [WARN ] i.g.h.e.GatlingHttpListener - Request 'Create Task' failed for user 2449
io.netty.channel.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
at io.netty.channel.nio.AbstractNioChannel$AbstractNioUnsafe$1.run(AbstractNioChannel.java:263)
at io.netty.util.concurrent.PromiseTask$RunnableAdapter.call(PromiseTask.java:38)
at io.netty.util.concurrent.ScheduledFutureTask.run(ScheduledFutureTask.java:127)
at io.netty.util.concurrent.AbstractEventExecutor.safeExecute(AbstractEventExecutor.java:163)
at io.netty.util.concurrent.SingleThreadEventExecutor.runAllTasks(SingleThreadEventExecutor.java:416)
at io.netty.channel.nio.NioEventLoop.run(NioEventLoop.java:515)
at io.netty.util.concurrent.SingleThreadEventExecutor$5.run(SingleThreadEventExecutor.java:918)
at io.netty.util.internal.ThreadExecutorMap$2.run(ThreadExecutorMap.java:74)
at io.netty.util.concurrent.FastThreadLocalRunnable.run(FastThreadLocalRunnable.java:30)
at java.lang.Thread.run(Thread.java:748)
14:38:44.397 [WARN ] i.g.h.e.r.DefaultStateProcessor - Request 'Create Task' failed for user 2449: i.n.c.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
14:38:44.398 [WARN ] i.g.h.e.GatlingHttpListener - Request 'Create Task' failed for user 2336
io.netty.channel.ConnectTimeoutException: connection timed out: backend/172.24.0.11:8080
at io.netty.channel.nio.AbstractNioChannel$AbstractNioUnsafe$1.run(AbstractNioChannel.java:263)
```

Here's the visualization of Spring Boot backend metrics, as we see there's lots of requests and mostly 2xx statuses.




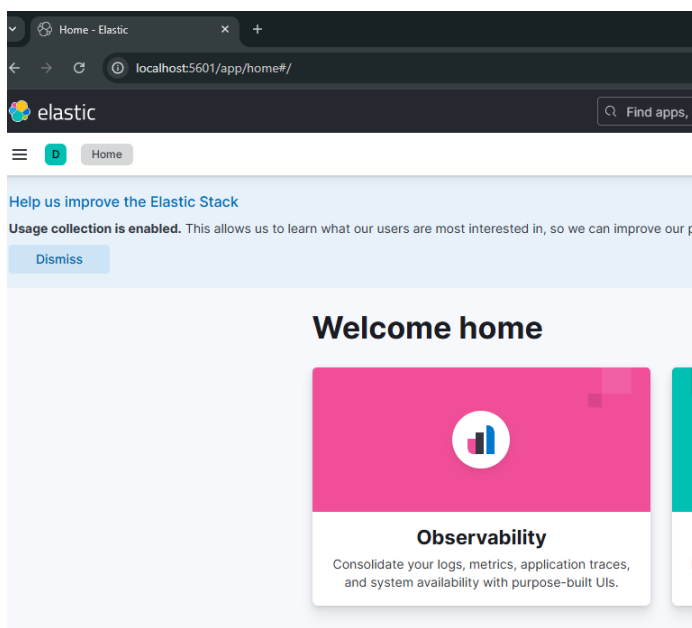
If we take a look at the JVM metrics now, we will see that the CPU and Memory utilization is quite high (or at least was quite high at its peak). The app has been running for the last 5 minutes and utilization is degrading as the requests are not loading the server anymore...



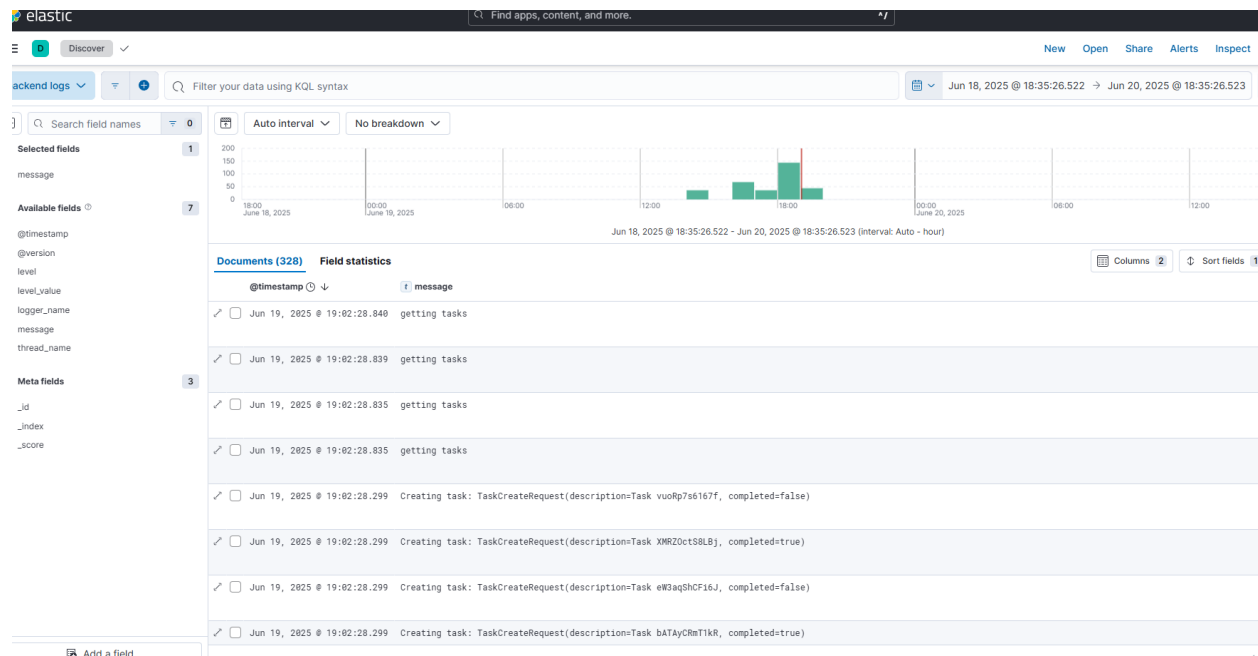


Logging

The Logging part is also interesting, where I defined ELK stack gathering logs from the backend to elastic with Logstash and visualizing them with Kibana. We need to go to the dashboards define the pattern and see the logs from there, currently my logs are under this elastic index as we see.

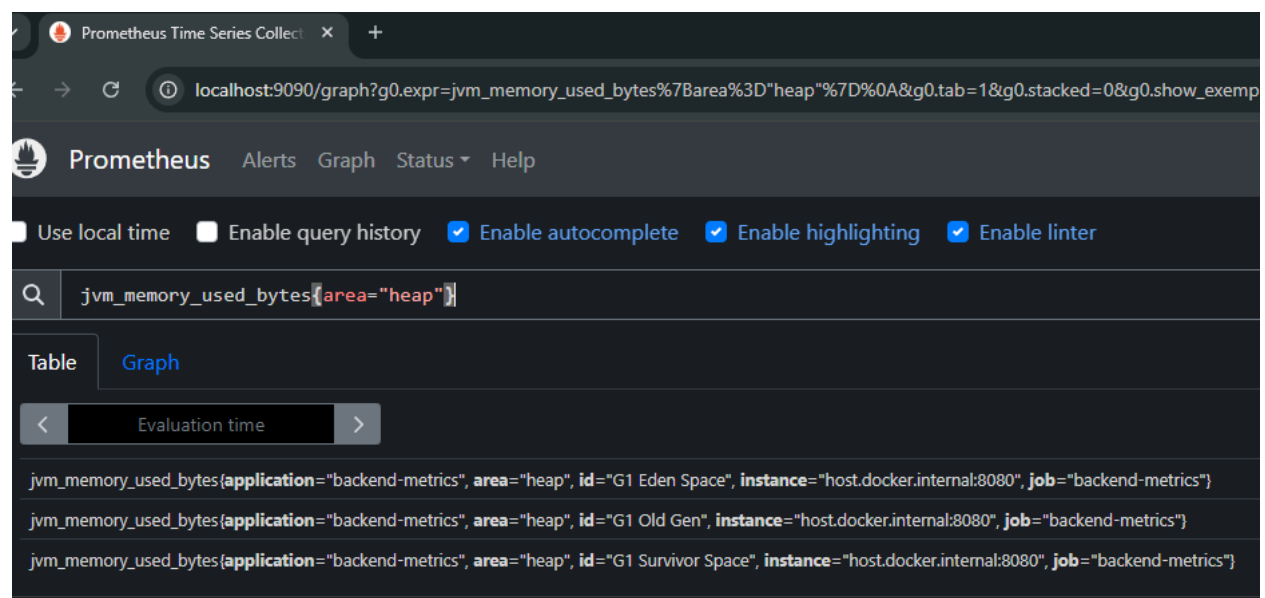


And if we go to discover page, select the message field from the left, we'll see the application logs as shown here.

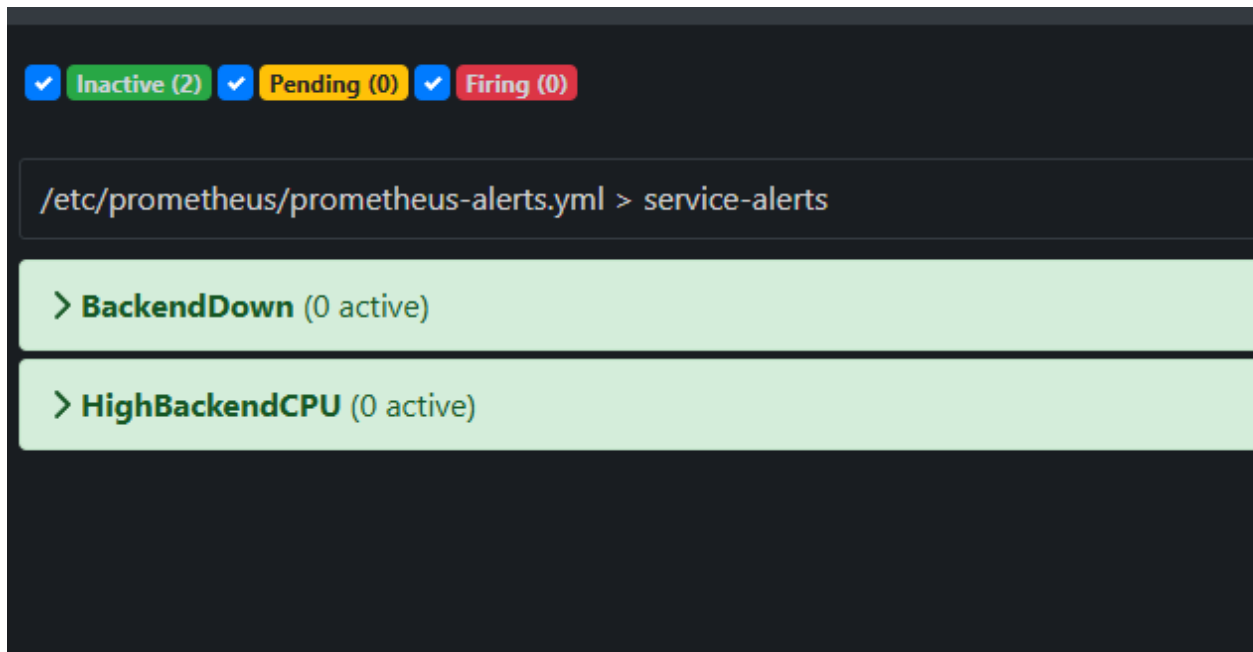


Prometheus

And here, Prometheus shows whatever metric query we give to it, in this case we see the consumption in the heap area of JVM selected by this query

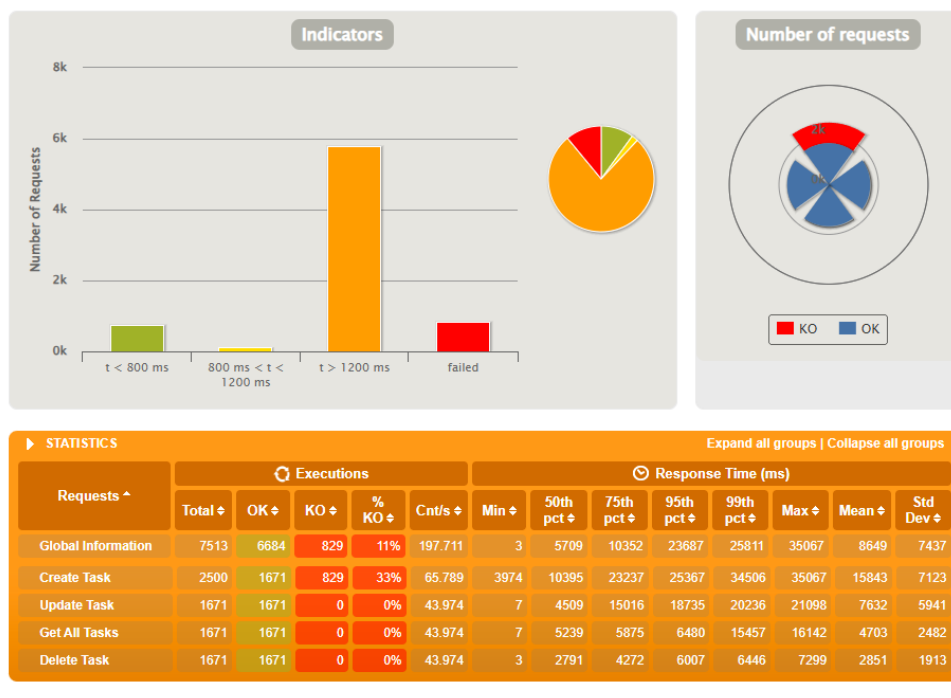


Alerting tasks which to we will come back later at **Incident Simulation**

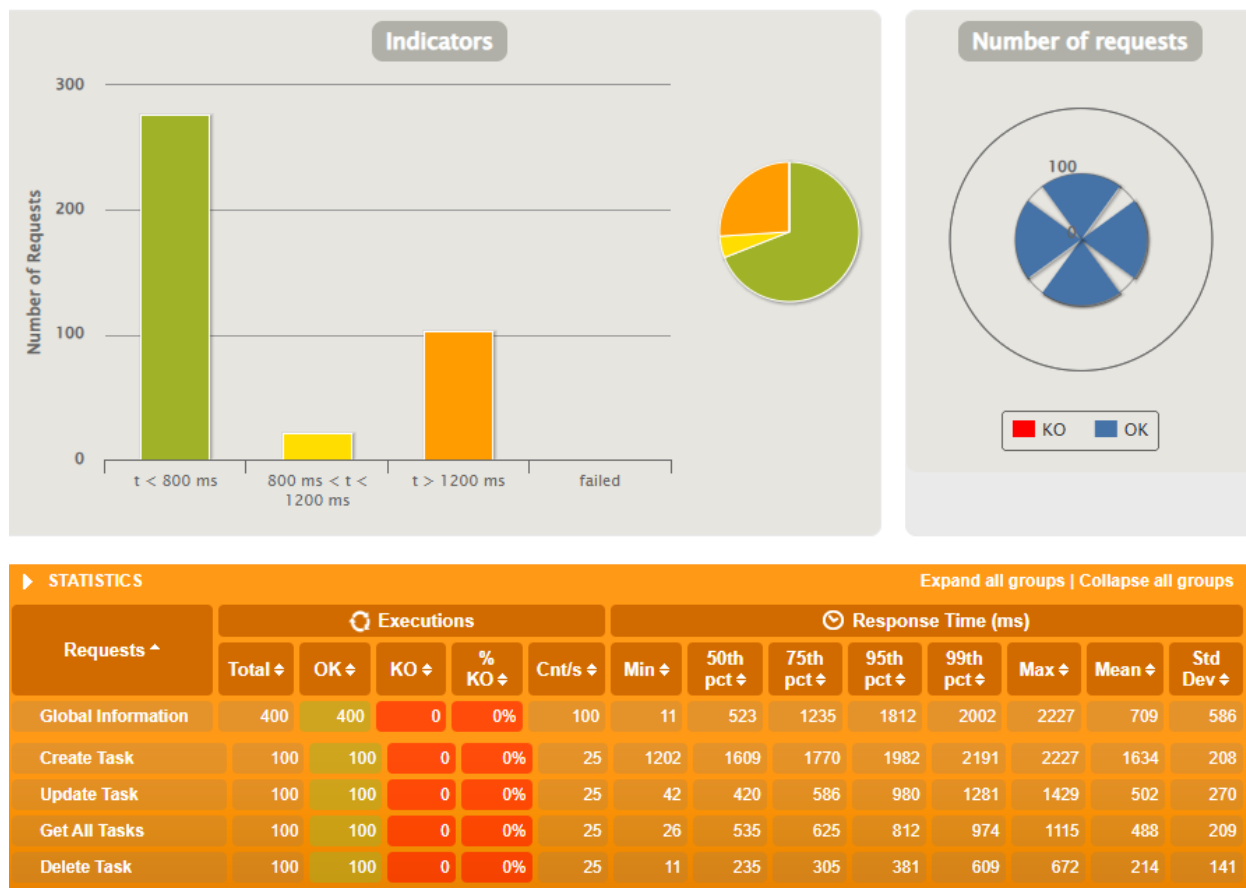


Gatling Simulation Results

2000+ Concurrent Requests



100 Concurrent Request



7. Security

```
scan-all.sh x
1  #!/bin/bash
2  set -e
3
4  IMAGES="
5  devops-final-prj-backend:latest
6  devops-final-prj-frontend:latest
7  nginx:1.25-alpine
8  prom/prometheus:v2.44.0
9  prom/alertmanager:v0.26.0
10 grafana/grafana:9.5.2
11 denvarez/gatling:latest
12 docker.elastic.co/elasticsearch/elasticsearch:8.13.4
13 docker.elastic.co/logstash/logstash:8.13.4
14 docker.elastic.co/kibana/kibana:8.13.4
15 "  Giorgi.Gagnidze, Yesterday · add trivy scan script
16
17 OUTPUT="trivy-report.txt"
18 # shellcheck disable=SC2188
19 > "$OUTPUT"
20
21 for image in $IMAGES; do
22     echo "===== " >> "$OUTPUT"
23     echo "Scanning $image ..." | tee -a "$OUTPUT"
24     trivy image "$image" | tee -a "$OUTPUT"
25 done
```

The requirement was to run **Trivy** scan on all Docker images, so I decided to write a shell script, in which there would be a hardcoded names of Docker images and it would scan everything automatically

```
Terminal Local - Ubuntu-20.04 (2) x + v
giorgi@giorgi:/mnt/c/Users/giorg/Desktop/devops-final-project/infra$ bash scan-all.sh
Scanning devops-final-prj-backend:latest ...
2025-06-19T19:16:18+04:00 WARN [vuln] Trivy DB may be corrupted and will be re-downloaded. If you manually downloaded DB - use the '--skip-db-update' flag to skip updating DB.
2025-06-19T19:16:18+04:00 INFO [vuln] Need to update DB
2025-06-19T19:16:18+04:00 INFO [vuln] Downloading vulnerability DB...
2025-06-19T19:16:18+04:00 INFO [vuln] Downloading artifact... repo="mirror.gcr.io/aquasec/trivy-db:2"
4.83 MiB / 65.88 MiB [----->] 7.33% 5.00 MiB p/s ETA 12s
```

This generated a huge file, or more than 20K lines, considering the fact that we have many images in our Docker compose.

Let's try to find and address any of the important issues that Trivy found. Most of the issues are version related, so upgrading package and image versions should help.

In my case for example, I can just upgrade Spring Boot version to avoid this error.

Target	Type	Vulnerabilities	Secrets
devops-final-prj-backend:latest (alpine 3.21.3)	alpine	0	-
app/backend-0.0.1.jar	jar	3	-

Legend:
- '-': Not scanned
- '0': Clean (no security findings detected)

Java (jar)
=====

Total: 3 (UNKNOWN: 0, LOW: 0, MEDIUM: 2, HIGH: 1, CRITICAL: 0)

Library	Vulnerability	Severity	Status	Installed Version	Fixed Version	Title
org.apache.tomcat.embed:tomcat-embed-core (backend-0.0.1.jar)	CVE-2025-48988	HIGH	fixed	10.1.41	11.0.8, 10.1.42, 9.0.106	tomcat: Apache Tomcat DoS in multipart upload https://avd.aquasec.com/nvd/cve-2025-48988
	CVE-2025-49125	MEDIUM				tomcat: Apache Tomcat: Security constraint bypass for pre/post-resources https://avd.aquasec.com/nvd/cve-2025-49125
org.springframework:spring-web (backend-0.0.1.jar)	CVE-2025-41234			6.2.7	6.2.8, 6.1.21	springframework: Reflected download attack in Spring Framework with non-ASCII headers https://avd.aquasec.com/nvd/cve-2025-41234

After upgrading the versions the vulnerabilities are gone as we see:

TerminalLocal × Ubuntu-20.04 (2) × + ▾

2025-06-19T19:53:32+04:00INFONumber of language-specific filesnum=1

2025-06-19T19:53:32+04:00INFO[jar] Detecting vulnerabilities...

Report Summary

Target	Type	Vulnerabilities	Secrets
devops-final-prj-backend:latest (alpine 3.21.3)	alpine	0	-
app/backend-0.0.1.jar	jar	0	-

Legend:
- '-': Not scanned
- '0': Clean (no security findings detected)

giorgi@Giorgi: /mnt/c/Users/giorg/Desktop/devops-final-project/infra\$

dependencies {

implementation 'org.springframework.boot:spring-boot-starter-web'

implementation 'org.springframework.boot:spring-boot-starter-data-jpa'

implementation 'org.springframework.boot:spring-boot-starter-actuator'

implementation 'org.springdoc:springdoc-openapi-starter-webmvc-ui:2.6.0'

}

262728293031

dependencies {

implementation 'org.apache.tomcat.embed:tomcat-embed-core:10.1.42'

implementation 'org.apache.tomcat.embed:tomcat-embed-websocket:10.1.42'

implementation 'org.apache.tomcat.embed:tomcat-embed-el:10.1.42'

implementation 'org.springframework:spring-web:6.2.8'

}

262728293031










This means we have successfully utilized Trivy to find vulnerabilities in our Docker images and fixed them accordingly.

Managing Sensitive Data

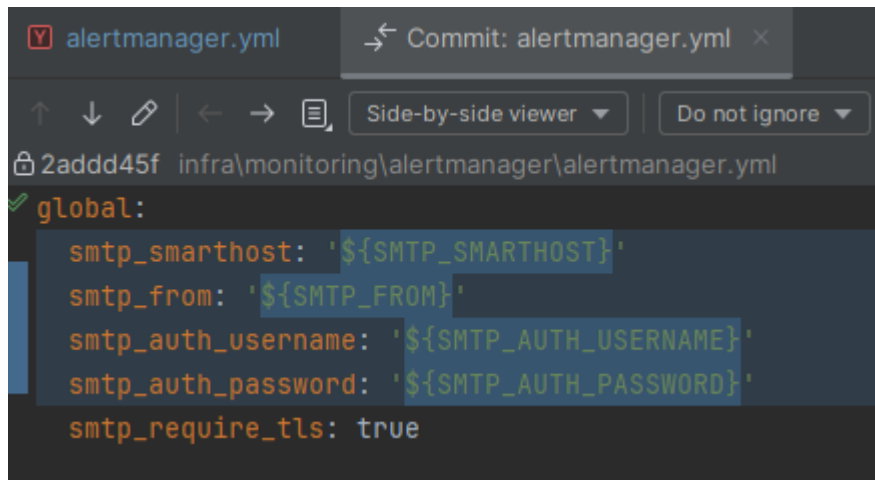
As you might have noticed, in my compose.yml I have specified the .env files for the images, for example the backend image has its own env file which is located at infra\backend\prod.env

```
services:
  backend:
    container_name: backend
    restart: always
    build:
      context: ../backend
      dockerfile: Dockerfile
    ports:
      - "8080:8080"
    depends_on:
      - grafana
      - prometheus
      - alertmanager
      - gatling
    env_file:
      - backend/prod.env
```

Then the upload of credentials are done via GitHub Actions using Ansible, copying the GitHub secrets and writing them to the VM we are deploying to.

 <code>SPRING_DATASOURCE_PASSWORD</code>	yesterday		
 <code>SPRING_DATASOURCE_URL</code>	yesterday		
 <code>SPRING_DATASOURCE_USERNAME</code>	yesterday		

Another example for this is Prometheus Alert Manager credentials

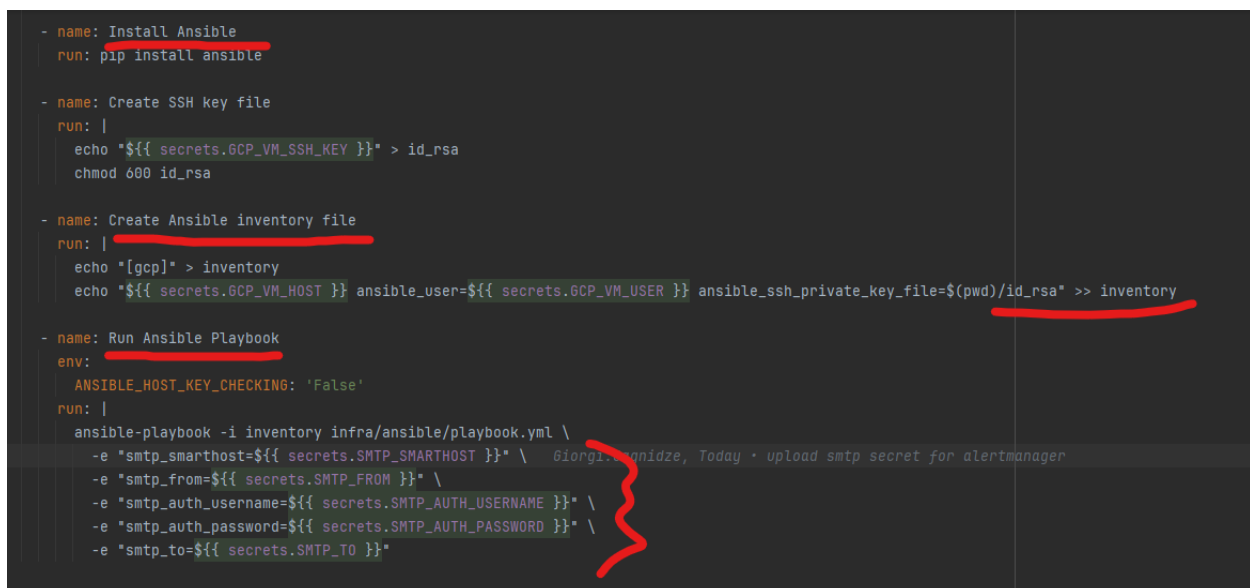


```
Y alertmanager.yml Commit: alertmanager.yml x
Side-by-side viewer Do not ignore
2add45f infra\monitoring\alertmanager\alertmanager.yml
global:
  smtp_smarthost: '${SMTP_SMARTHOST}'
  smtp_from: '${SMTP_FROM}'
  smtp_auth_username: '${SMTP_AUTH_USERNAME}'
  smtp_auth_password: '${SMTP_AUTH_PASSWORD}'
  smtp_require_tls: true
```

This is also managed by Ansible and GitHub actions secrets, this alertmanager.yml is so called env file for the alert manager.

As I already mentioned, infrastructure is automatically deployed by infra-deploy.yml in GitHub Actions, which also runs Ansible automatically and creates inventory file for it.

It passes secret arguments and Ansible deploys alertmanager.yml file on the server itself, without leaking the credentials.



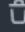














```
- name: Install Ansible
  run: pip install ansible

- name: Create SSH key file
  run: |
    echo "${{ secrets.GCP_VM_SSH_KEY }}" > id_rsa
    chmod 600 id_rsa

- name: Create Ansible inventory file
  run: |
    echo "[gcp]" > inventory
    echo "${{ secrets.GCP_VM_HOST }} ansible_user=${{ secrets.GCP_VM_USER }} ansible_ssh_private_key_file=$(pwd)/id_rsa" >> inventory

- name: Run Ansible Playbook
  env:
    ANSIBLE_HOST_KEY_CHECKING: 'False'
  run: |
    ansible-playbook -i inventory infra/ansible/playbook.yml \
      -e "smtp_smarthost=${{ secrets.SMTP_SMARTHOST }}" \
      -e "smtp_from=${{ secrets.SMTP_FROM }}" \
      -e "smtp_auth_username=${{ secrets.SMTP_AUTH_USERNAME }}" \
      -e "smtp_auth_password=${{ secrets.SMTP_AUTH_PASSWORD }}" \
      -e "smtp_to=${{ secrets.SMTP_TO }}"
```

 SMTP_AUTH_PASSWORD	7 hours ago		
 SMTP_AUTH_USERNAME	7 hours ago		
 SMTP_FROM	7 hours ago		
 SMTP_SMARTHOST	7 hours ago		
 SMTP_TO	7 hours ago		

Therefore, the credentials and secrets are managed by .env files, GitHub Secrets and Ansible.

8. Incident Simulation + Post-Mortem

Post mortem file and discussions is attached with related screenshots.

For the incident simulation part I have taken down backend container manually, another option could've been stress loading backend with Gatling until it would shut down automatically, and could've had automatic restart policy for Docker but that would also kill my PC. By default I have restart: auto so if backend went down on itself it would've restarted

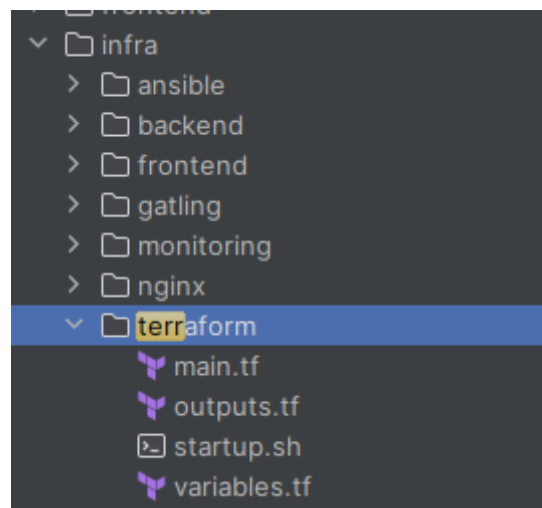
9. Automation (Ansible, Terraform, GitHub Actions)

As per bonus, it was suggested to use **Ansible**, which I have used in many ways. As I've discussed with the head of the class (course), I have deploy the entire application to the cloud environment, on GCP VM.

I have provisioned the infrastructure with **Terraform** and automated deployment with **Ansible** and **GitHub Actions**

Terraform

As I have already mentioned, under my /infra is everything that the local/VM infrastructure needs. Under Terraform folder I have defined the script for provisioning an instance on GCP and allowing multi-port traffic on this VM. Also I have multiple variables here, such as VM's name, machine type, Region, Zone, Service Account Key, etc...



```
provider "google" {  
  credentials = file(var.gcp_credentials_file)  
  project     = var.project_id  
  region      = var.region  
  zone        = var.zone  
}  
  
resource "google_compute_instance" "app" {  
  name         = var.vm_name  
  machine_type = var.machine_type  
  zone         = var.zone  
  
  boot_disk {  
    initialize_params {  
      image = var.disk_image  
    }  
  }  
}
```

This main.tf is run by GitHub Actions workflow, terraform.yml, which only runs whenever something changes under the following path “infra/terraform/” or it can also be run manually from GitHub. Here are the steps, they’re pretty straightforward:

```
- name: Set up Terraform
  uses: hashicorp/setup-terraform@v3

- name: Write GCP credentials file
  run: |
    echo '${{ secrets.GCP_CREDENTIALS }}' > gcp-key.json

- name: Write SSH pubkey
  run: |
    echo '${{ secrets.TF_SSH_PUB_KEY }}' > id_rsa.pub

- name: Terraform Init
  run: terraform init

- name: Terraform Apply
  run: |
    terraform apply -auto-approve \
      -var="gcp_credentials_file=gcp-key.json" \
      -var="project_id=${{ secrets.GCP_PROJECT_ID }}" \
      -var="region=europe-west3" \
      -var="zone=europe-west3-a" \
      -var="ssh_user=deployer" \
      -var="ssh_pub_key=id_rsa.pub"

- name: Show Output IP
  run: terraform output instance_ip
```

As you see, most of the credentials are taken from GitHub secrets. A successful run for this action can be seen on my repository URL, under the actions section.

Provision infrastructure on GCP with Terraform
terraform.yml

4 workflow runs

This workflow has a workflow_dispatch event trigger. [Run workflow](#)

	Event	Status	Branch	Actor
✓ add instance User data (init script)	Provision infrastructure on GCP with Terraform #4: Commit 2410a0d pushed by giorgigagnidze16	main	yesterday	33s
✗ rename job	Provision infrastructure on GCP with Terraform #3: Commit 3dc787d pushed by giorgigagnidze16	main	yesterday	9s


```
129
130 Plan: 2 to add, 0 to change, 0 to destroy.
131
132 Changes to Outputs:
133   + instance_ip = (known after apply)
134 google_compute_firewall.default-allow-ssh-http-https: Creating...
135 google_compute_instance.app: Creating...
136 google_compute_firewall.default-allow-ssh-http-https: Still creating... [00m10s elapsed]
137 google_compute_instance.app: Still creating... [00m10s elapsed]
138 google_compute_firewall.default-allow-ssh-http-https: Creation complete after 12s [id-projects/***/global/firewalls/default-allow-ssh-http-https]
139 google_compute_instance.app: Creation complete after 19s [id-projects/***/zones/europe-west3-a/instances/devops-vm]
140
141 Apply complete! Resources: 2 added, 0 changed, 0 destroyed.
142
143 Outputs:
144
145 instance_ip = "34.40.124.152"
```

▼

✓ Show Output IP

0s

```
1 ▶ Run terraform output instance_ip
6 "34.40.124.152"
```

And there is my instance, the IP in the both images are different since the public IP is not static in this case and changes after every stop-start.

VM instances [Create instance](#) [Import VM](#) [Refresh](#) [Learn](#)

[Instances](#) [Observability](#) [Instance schedules](#)

VM instances

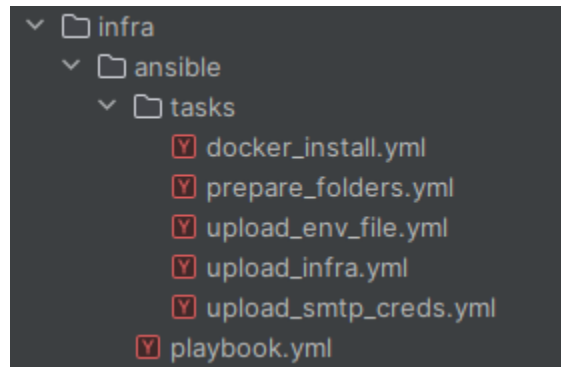
Filter Enter property name or value

<input type="checkbox"/> Status	Name ↑	Zone	Recommendations	In use by	Internal IP	External IP	Connect
<input type="checkbox"/> ✓	devops-vm	europe-west3-a			<div></div> (nic0)	34.40.97.161 (nic0)	SSH ▾ ⋮

Now, we have seen I have used Terraform to provision the infra on GCP, Let's get back to Ansible and how I have used it.

Ansible

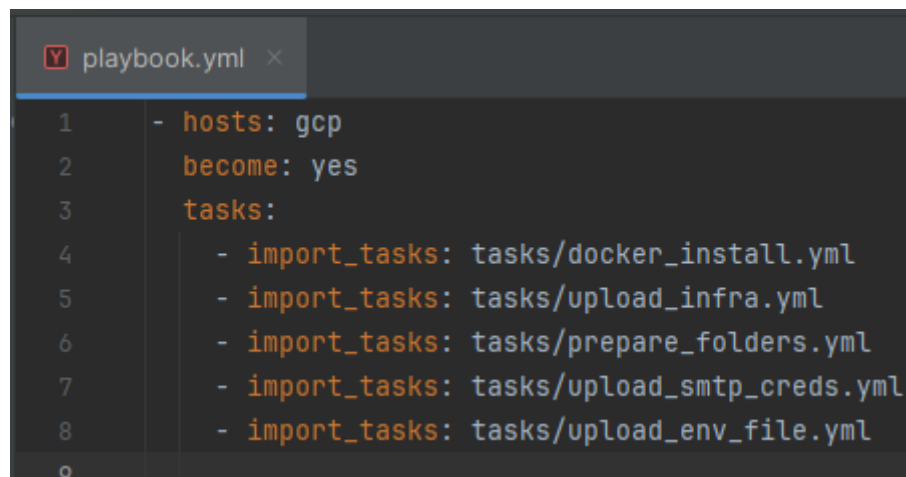
I have excluded Terraform folder from infra deployment on cloud, since the tf files are not really needed on the remote VM, we only need to provision Docker, make sure deployment folders exists and so on. That's what I have used Ansible for.



Under Ansible folder, I have main playbook.yml file for running Ansible tasks.

Tasks are for:

- Docker installation
- Preparing deployment folders
- Upload infra folder with compose to the VM
- Upload env files and credentials to the VM with help of GitHub secrets



Let's skip Docker install part, since its most complicated one out of the 3 and more unrelated to our current domain, its Debian related.

Upload Infra:

Uploads everything under /infra folder to remote host, excluding terraform, ansible and .git

My GCP deployment folder is /tmp/deployment

```
upload_infra.yml x
1  - name: Ensure rsync is installed
2    ansible.builtin.apt:
3      name: rsync
4      state: present
5      update_cache: yes
6
7  - name: delete infra directory
8    ansible.builtin.file:
9      path: /tmp/deployment/infra
10     state: absent
11
12  - name: recreate an empty infra directory  Giorgi.Gagnidze,
13    ansible.builtin.file:
14      path: /tmp/deployment/infra
15      state: directory
16      mode: '0755'
17      owner: "{{ ansible_user }}"
18
19  - name: Copy all folders except terraform
20    ansible.builtin.synchronize:
21      src: "{{ playbook_dir }}/../"
22      dest: /tmp/deployment/infra/
23      rsync_opts:
24        - "--exclude=terraform"
25        - "--exclude=.git"
26        - "--exclude=ansible"
```

Prepare Folders:

This ensures that the backend and frontend directories exist, upon deployment the directories are emptied and created again to avoid any inconsistencies between versions and etc...

```
❏ prepare_folders.yml ×
1  - name: delete backend directory
2    ansible.builtin.file:
3      path: /tmp/deployment/backend
4      state: absent
5  |   Giorgi.Gagnidze, Today • create deployment folders on
6  - name: recreate an empty backend directory
7    ansible.builtin.file:
8      path: /tmp/deployment/backend
9      state: directory
10     mode: '0755'
11     owner: "{{ ansible_user }}"
12
13  - name: delete frontend directory
14    ansible.builtin.file:
15      path: /tmp/deployment/frontend
16      state: absent
17
18  - name: recreate an empty frontend directory
19    ansible.builtin.file:
20      path: /tmp/deployment/frontend
21      state: directory
22      mode: '0755'
23      owner: "{{ ansible_user }}"
```

Upload SMTP Creds:

Uploads alertmanager.yml to the remote host after GitHub Actions passes the arguments down to the playbook.yml

```
upload_smtp_creds.yml x
1  - name: upload alert manager creds env file
2    ansible.builtin.copy:
3      dest: /tmp/deployment/infra/monitoring/alertmanager/alertmanager
4      content: |
5        global:
6          smtp_smarthost: '{{ smtp_smarthost }}'
7          smtp_from: '{{ smtp_from }}'   Giorgi.Gagnidze, Today • upl
8          smtp_auth_username: '{{ smtp_auth_username }}'
9          smtp_auth_password: '{{ smtp_auth_password }}'
10         smtp_require_tls: true
11
12        route:
13          receiver: 'email-alerts'
14
15        receivers:
16          - name: 'email-alerts'
17            email_configs:
18              - to: '{{ smtp_to }}'
19                send_resolved: true
20        owner: "{{ ansible_user | default('deployer') }}"
21        mode: '0600'
```

GitHub Actions + Ansible

Infra-deploy.yml is triggered when a new change is committed under infra/ansible/ directory. Ansible is installed on the runner, GCP key is written on the id_rsa file, passed down to inventory and finally Ansible Playbook is run with provided secret arguments.

```
python-version: 3.12

- name: Install Ansible
  run: pip install ansible

- name: Create SSH key file
  run: |
    echo "${{ secrets.GCP_VM_SSH_KEY }}" > id_rsa
    chmod 600 id_rsa

- name: Create Ansible inventory file
  run: |
    echo "[gcp]" > inventory
    echo "${{ secrets.GCP_VM_HOST }}" ansible_user=${{ secrets.GCP_VM_USER }} ansible_ssh_private_key_file=$(pwd) >> inventory

- name: Run Ansible Playbook
  env:
    ANSIBLE_HOST_KEY_CHECKING: 'False'
  run: |
    ansible-playbook -i inventory infra/ansible/playbook.yml \
      -e "smtp_smarthost=${{ secrets.SMTP_SMARTHOST }}" \
      -e "smtp_from=${{ secrets.SMTP_FROM }}" \
      -e "smtp_auth_username=${{ secrets.SMTP_AUTH_USERNAME }}" \
      -e "smtp_auth_password=${{ secrets.SMTP_AUTH_PASSWORD }}" \
      -e "smtp_to=${{ secrets.SMTP_TO }}"
```

Backend Deploy

There's an entire pipeline for backend deploy, Unit tests need to pass, linting checks and building jar. After that, remote backend deployment directory is emptied, Dockerfile uploaded as well as built jar and backend container is started

```
backend.yml x
9  jobs:
10    build-test-upload:
13      steps:
38
39      - name: empty deployment directory
40        uses: appleboy/ssh-action@v1.0.0
41        with:
42          host: ${ secrets.GCP_VM_HOST }
43          username: ${ secrets.GCP_VM_USER }
44          key: ${ secrets.GCP_VM_SSH_KEY }
45          script: |
46            rm -rf /tmp/deployment/backend/*
47
48      - name: upload Dockerfile
49        uses: appleboy/scp-action@v0.1.4
50        with:
51          host: ${ secrets.GCP_VM_HOST }
52          username: ${ secrets.GCP_VM_USER }
53          key: ${ secrets.GCP_VM_SSH_KEY }
54          source: ./backend/Dockerfile
55          target: /tmp/deployment/
56      | Giorgi.Gagnidze, Yesterday · backend ci
57      - name: upload build/libs
58        uses: appleboy/scp-action@v0.1.4
59        with:
60          host: ${ secrets.GCP_VM_HOST }
61          username: ${ secrets.GCP_VM_USER }
62          key: ${ secrets.GCP_VM_SSH_KEY }
63          source: ./backend/build/libs
64          target: /tmp/deployment/
65
66      - name: start backend container
67        uses: appleboy/ssh-action@v1.0.0
68        with:
69          host: ${ secrets.GCP_VM_HOST }
70          username: ${ secrets.GCP_VM_USER }
71          key: ${ secrets.GCP_VM_SSH_KEY }
72          script: |
73            cd /tmp/deployment/infra
74            docker compose build --no-cache backend
75            docker compose up -d backend
```

Frontend Deploy

Front end folder uploaded to the remote host, image is build and container started!

```
- name: upload frontend
  uses: appleboy/scp-action@v0.1.4
  with:
    host: ${ secrets.GCP_VM_HOST }
    username: ${ secrets.GCP_VM_USER }
    key: ${ secrets.GCP_VM_SSH_KEY }
    source: ./frontend/
    target: /tmp/deployment/

- name: Start frontend container
  uses: appleboy/ssh-action@v1.0.0
  with:
    host: ${ secrets.GCP_VM_HOST }
    username: ${ secrets.GCP_VM_USER }
    key: ${ secrets.GCP_VM_SSH_KEY }
    script: |
      cd /tmp/deployment/infra
      docker compose build --no-cache frontend
      docker compose up -d frontend
```

In compose file itself, backend and frontend are both dependent on other services, and their startup also triggers other services to start.

Backend Deployment:

This workflow has a <code>workflow_dispatch</code> event trigger.				Run workflow ▾
✓	add elk monitoring	main	7 hours ago 2m 58s	...
Backend CI & Deploy #20: Commit <code>2add45</code> pushed by giorgigagnidze16				
✓	Backend CI & Deploy	main	9 hours ago 2m 40s	...
Backend CI & Deploy #19: Manually run by giorgigagnidze16				

Frontend Deployment:

4 workflow runs		Event ▾	Status ▾	Branch ▾	Actor ▾
This workflow has a <code>workflow_dispatch</code> event trigger.				Run workflow ▾	
✓	Frontend CI & Deploy	main	9 hours ago 4m 5s	...	
Frontend CI & Deploy #4: Manually run by giorgigagnidze16					
✗	Frontend CI & Deploy	main	9 hours ago 2m 12s	...	
Frontend CI & Deploy #3: Manually run by giorgigagnidze16					

The remote host machine after deploying infra using Ansible: As we see all the infra folder is uploaded

```
ssh.cloud.google.com/v2/ssh/projects/sincere-elixir-463316-t5/zones/e

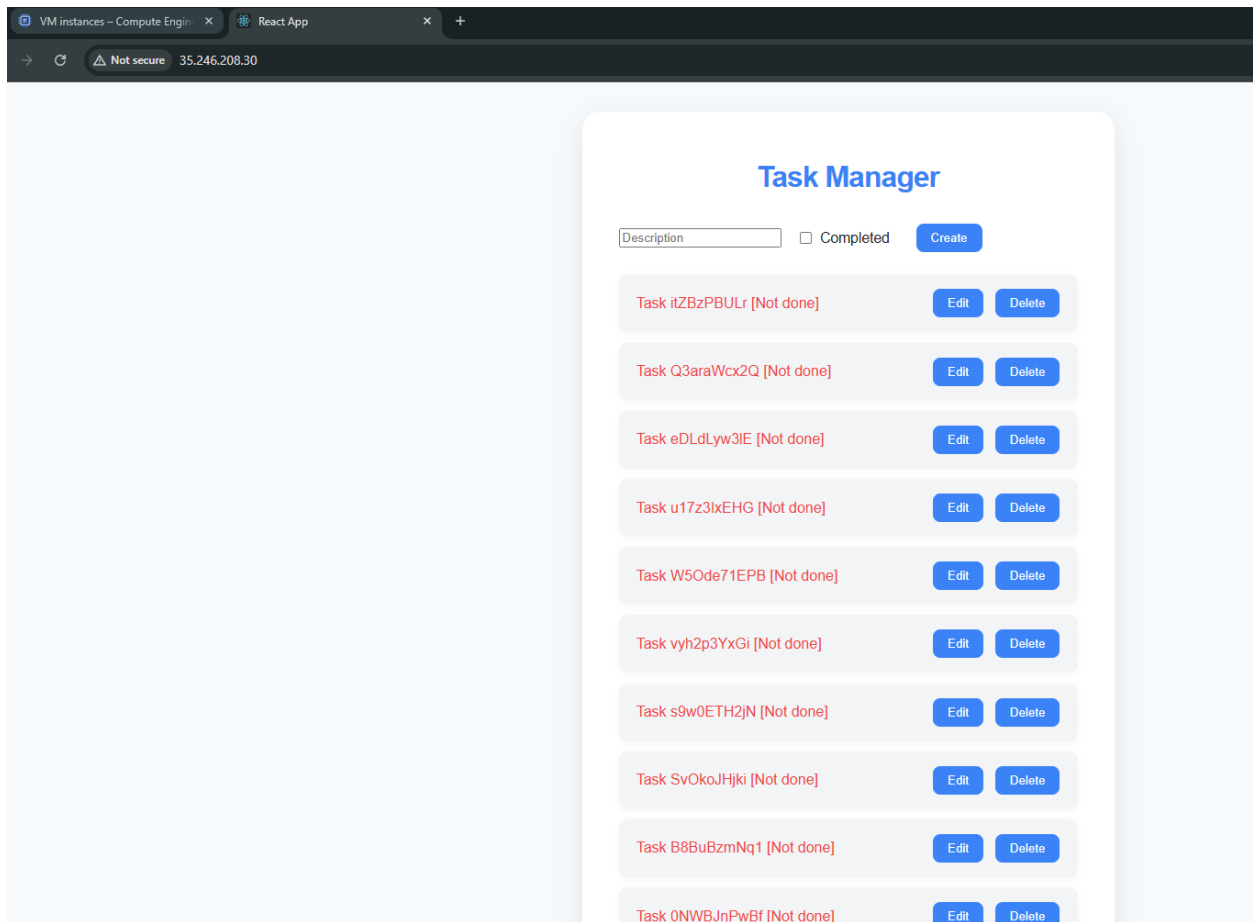
SSH-in-browser

root@devops-vm:/tmp/deployment/infra# tree
.
├── backend
│   └── prod.env
├── compose.yml
├── frontend
├── gatling
│   └── user-files
│       └── simulations
│           ├── BasicSimulation.scala
│           └── CrudSimulation.scala
├── monitoring
│   ├── alertmanager
│   │   └── alertmanager.yml
│   ├── elk
│   │   └── logstash.conf
│   ├── grafana
│   │   ├── dashboards
│   │   │   ├── 17175_rev2.json
│   │   │   └── 4701_rev10.json
│   │   └── provisioning
│   │       ├── dashboards
│   │       │   └── dashboard.yml
│   │       └── datasources
│   │           └── datasources.yml
│   ├── prometheus-alerts.yml
│   └── prometheus.yml
├── nginx
│   └── nginx.conf
├── scan-all.sh
└── trivy-report.txt

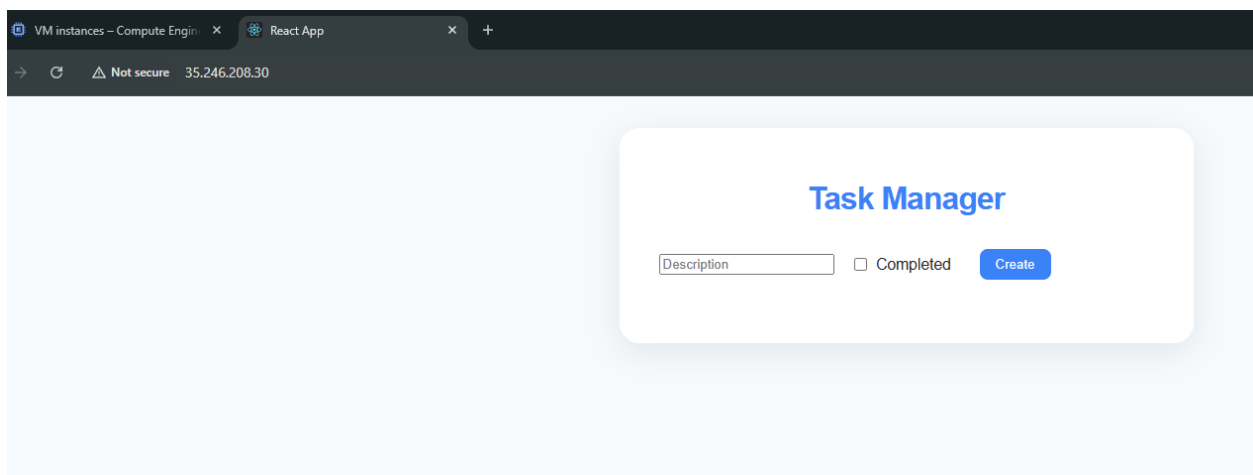
15 directories, 15 files
root@devops-vm:/tmp/deployment/infra#
```

```
root@devops-vm:/tmp/deployment# ls
backend  frontend  infra
root@devops-vm:/tmp/deployment# ls backend/
Dockerfile  build
root@devops-vm:/tmp/deployment# ls frontend/
Dockerfile  README.md  build  node_modules  package-lock.json  package.json  public  src  tsconfig.json
root@devops-vm:/tmp/deployment#
```

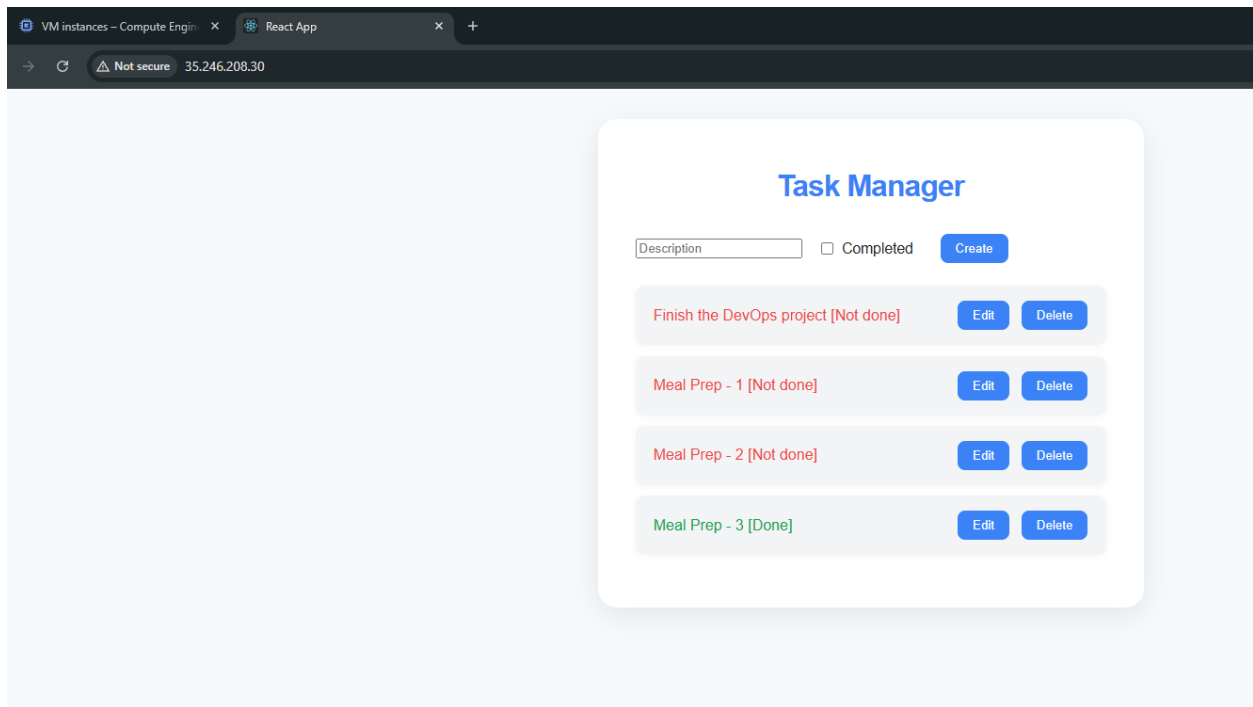
Let's make sure instance is running all the services smoothly. Live Demo:



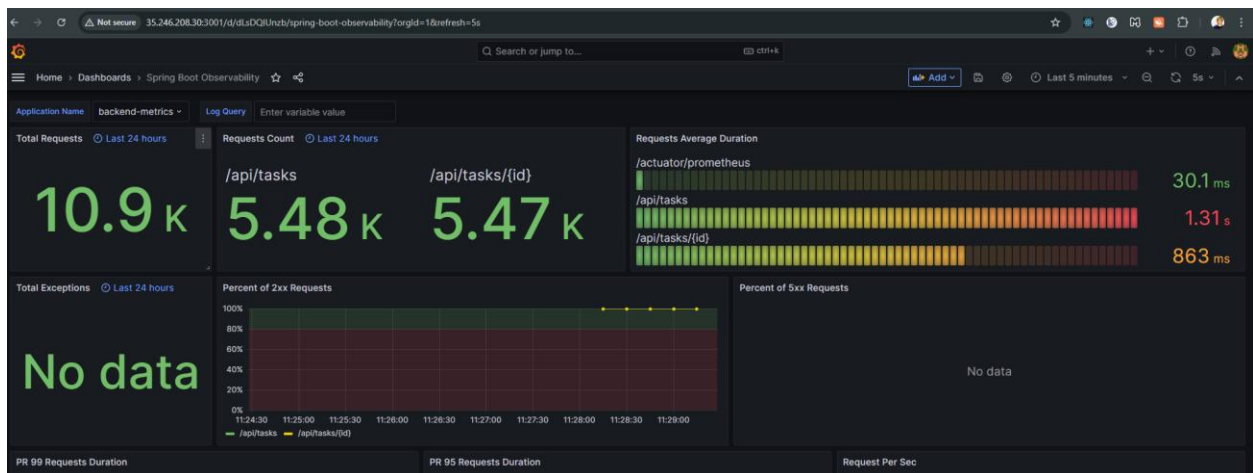
As we see these are tasks created by Gatling simulation, I was lucky I had the opportunity to see this actually since I simulated full CRUD and after deletions, we cant see the tasks obviously.



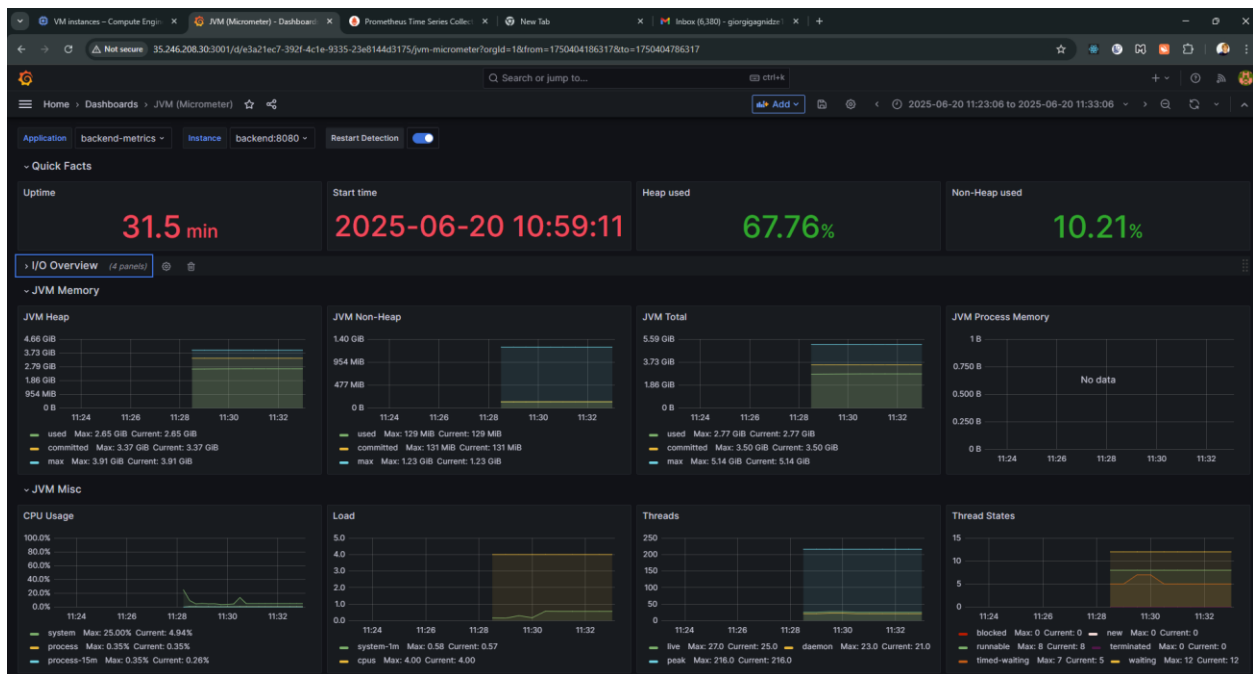
After refresh, they're gone and we can see that this thing is fully functional too.



We can also check Grafana, ELK and everything

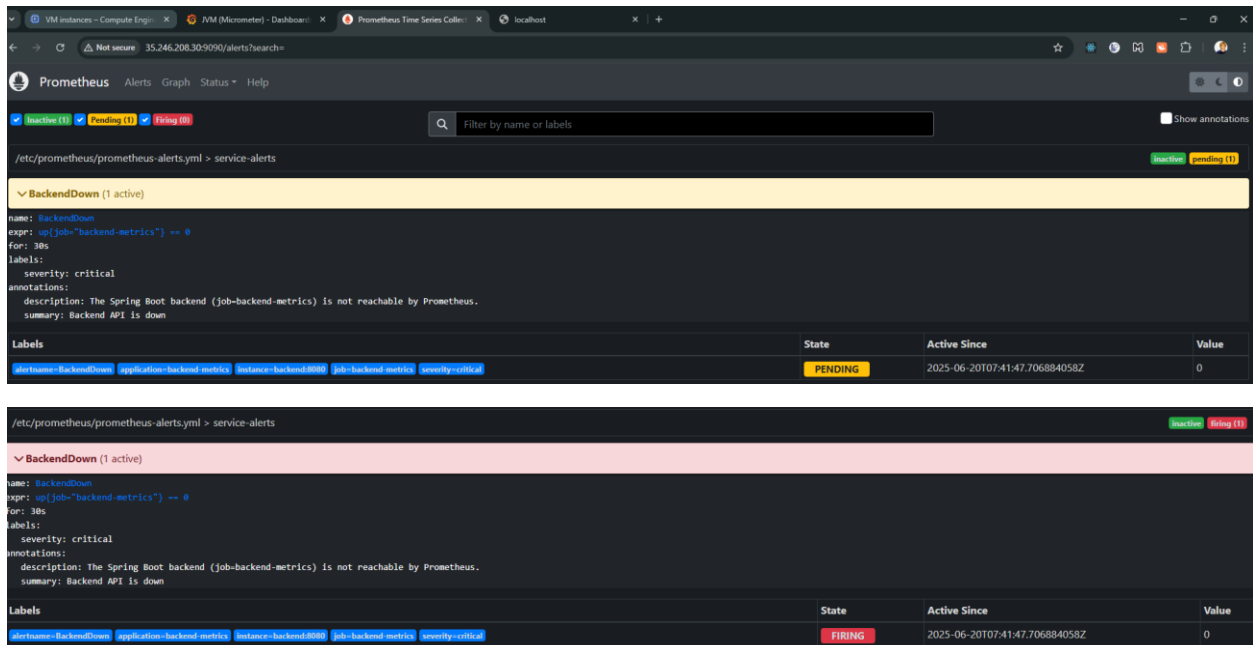


As you see there's a lots of requests from Gatling due to me restarting Gatling manually.



Ok, now let's try to take down the backend service and we'll see that Prometheus responds accordingly

```
prometheus
root@devops-vm:/tmp/deployment/infra/monitoring# docker compose down backend
[+] Running 2/2
 ✓ Container backend                               Removed
 ! Network devops-final-prj_default Resource is still in use
root@devops-vm:/tmp/deployment/infra/monitoring#
```



Well, that's it. I have utilized every tool described in the task, applying the best practices possible. Additionally did the complete CI/CD deployment of an application to the GCP cloud for which I used Terraform, to get an instance VM, which would handle many Docker images, and load that Gatling was putting on the backend as well as 16GB of ram. Additionally I have used ELK & Prometheus Alerts for monitoring, GitHub Actions for automated deployment and Ansible for infrastructure provisioning (package installations, folder creation, updates, etc...)